

COMP4650/6490 Document Analysis – Semester 2 / 2022

Assignment 1

Due 17:00 on Wednesday 17 August 2022 AEST (UTC +10)

Last updated July 25, 2022

Overview

In this assignment, your task is to index a document collection into an inverted index, and then measure search performance based on predefined queries. A document collection containing more than 30,000 government site descriptions is provided for this assignment, along with a set of queries (gov/topics/gov.topics) and the expected returned documents (gov/qrels/gov.qrels). The provided code implements most of an information retrieval system. This provided code is designed to be simple to understand and modify, it is not efficient nor scalable. When developing a real-world IR system you would be better off using high performance software such as Apache Lucene¹.

Throughout this assignment:

1. You will develop a better understanding of indexing, including the tokenizer, parser, and normaliser components, and how to improve the search performance given a predefined evaluation metric;
2. You will develop a better understanding of search algorithms, and how to obtain better search results, and
3. You will find the best performance.

Submission

- The answers to this assignment (including your modified code) will be submitted online in Wattle, see the link **Assignment 1 Submission** in Week 4 (15 to 19 August).
- You can edit your answers many times and they will be saved by Wattle.
- Note that **Wattle does not allow us to access any earlier edited versions of your answers, so check very carefully what you submit as the final version.**
- **You can only submit your assignment once. Make sure you submit the final version of your assignment answers before the submission deadline.**

Marking

This assignment will be marked out of 15, and it will contribute 15% of your final course mark.

Your answers to coding questions will be marked based on the quality of your code (is it efficient, is it readable, is it extendable, is it correct) and the solution in general (is it appropriate, is it reliable, does it demonstrate a suitable level of understanding).

Your answers to discussion questions will be marked based on how convincing your explanations are (are they sufficiently detailed, are they well-reasoned, are they backed by appropriate evidence, are they clear, do they use appropriate aids such as tables and plots where necessary).

¹<https://lucene.apache.org/>

Question 1: Implement the Indexer and Evaluate (3 marks)

Your first task is to implement index construction. You should complete the function stub `index_from_tokens` in `indexer.py`. This function takes as input a sorted list of tuples of the form `(token_string, doc_id)`, indicating that the token `token_string` was in the document represented by `doc_id`. For each occurrence of a token in a document there is a tuple in this list – this means duplicate tuples indicate term count. **The input list is sorted in ascending order by token_string then doc_id.**

Once complete the `index_from_tokens` function should output two dictionaries, the first dictionary `index` is a mapping from token strings to a sorted list of `(doc_id, term_frequency)` tuples. These lists should have their elements sorted in ascending order by `doc_id`, and contain only unique `doc_ids` – duplicates are used to compute term frequency. The second dictionary `doc_freq` is a mapping from `token_string` to document frequency.

Once you have implemented `index_from_tokens` you should run `indexer.py` to store the index, then run `query.py` to run a set of test queries, finally run `evaluate.py` to evaluate the query results against the ground truth. Record your evaluation results in your answers and make sure you submit your `indexer.py`.

Example Input for `index_from_tokens`:

```
[("cat", 1), ("cat", 1), ("cat", 2), ("door", 1), ("water", 3)]
```

Example Output for `index_from_tokens`:

```
index: {"cat": [(1, 2), (2, 1)], "door": [(1, 1)], "water": [(3, 1)]}
```

```
doc_freq: {"cat": 2, "door": 1, "water": 1}
```

Question 2: Imple

Currently `query_tfidf` in `query.py` but you will change it to use TF-IDF. In

your solution both the query and the document vectors should be TF-IDF vectors.

modify the `run_query` function and the `get_document_vector` function in `query_tfidf.py`.

The TF-IDF variant you should implement is:

$$\text{TF-IDF} = n_t \ln \frac{N}{1 + n_d}$$

Where n_t is the term frequency, n_d is the document frequency, and N is the total number of documents in the collection. This is almost the standard TF-IDF variant, except that 1 is added to the document frequency to avoid division by zero errors.

Once you have implemented TF-IDF cosine similarity, run the `query_tfidf.py` file, then run `evaluate.py`, and record the results in your answers. Make sure you submit your `query_tfidf.py`.

Question 3: Explore Linguistic Processing Techniques (3 marks)

For this question you will explore ways to improve the `process_tokens` function in `string_processing.py`. The current function removes stopwords. You should modify the function and explore the results. To modify the function, you should make changes to the functions `process_token_1`, `process_token_2`, and `process_token_3` and then uncomment the one you want to test within the main `process_tokens` function. You should pick **at least three different modifications** and evaluate them (you can add new process tokens functions if you want to evaluate more than three modifications). See lectures for some possible modifications. You might find the Python `nltk` library useful. The modifications you make do not need to require significant coding, the focus of this question is choosing reasonable modifications and explaining the results.

Query 1: Australasia OR logistic

Query 2: heart AND warm

Query 3: global AND space AND wildlife

Query 4: engine OR origin AND record AND wireless

Query 5: placement AND sensor OR max AND speed

Make sure you submit your code for this question (i.e. `query_boolean.py`) as well as your answers.

Question 5: Evaluating IR Systems (3 marks)

Answer the following two questions:

- (a) Explain how you can evaluate your boolean query system.
- (b) Now consider the case where the output of your boolean query system is going to be re-ranked and only the top 10 results displayed. Explain how you could evaluate this combined boolean query and re-ranking system.

Your answers to both (a) and (b) should at a minimum consider:

- What data you would need (e.g. queries and ground truth);
- The challenges that you would have to face getting this data;
- What metrics are appropriate, and why these metrics are appropriate

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro